



Call-normalised deliveries by prompt configuration and scenario



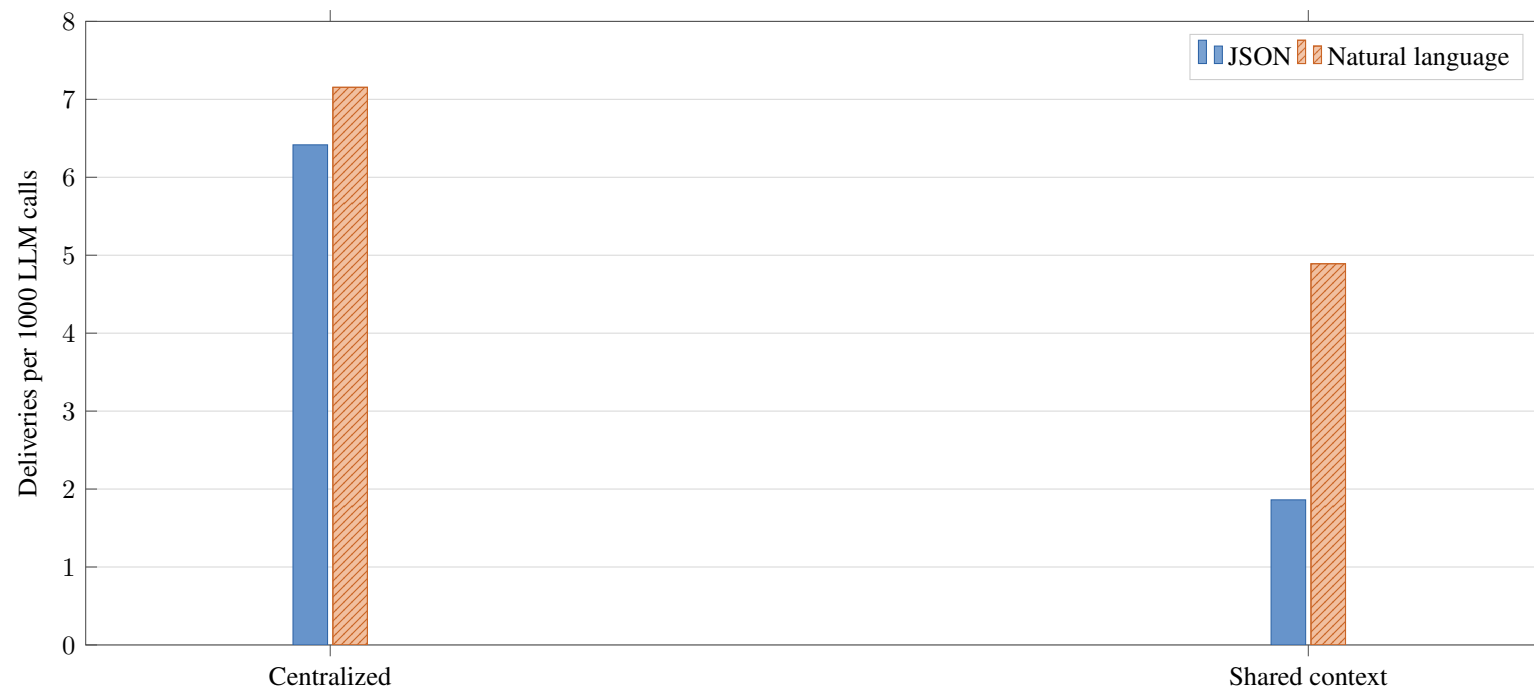
Six discrete scenario groups; 24 matched configuration-level observations per prompt configuration within each scenario.

Call-normalised deliveries by prompt configuration and model



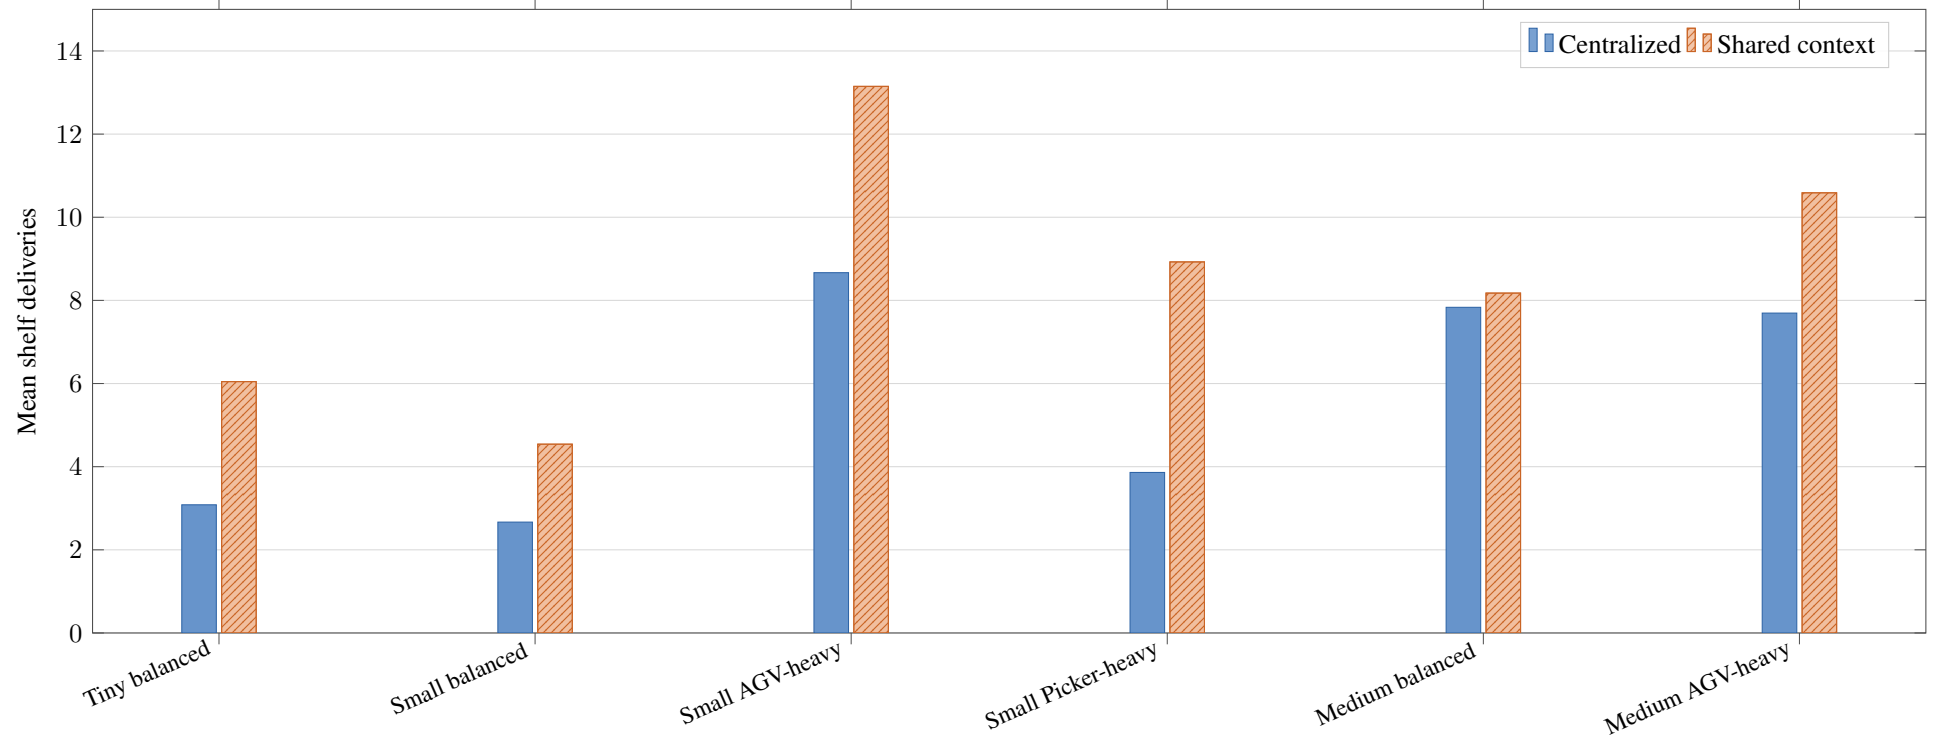
Models are ordered by nominal parameter size. Bars show discrete model means, with no interpolation between model families.

Prompt configuration within each planning architecture



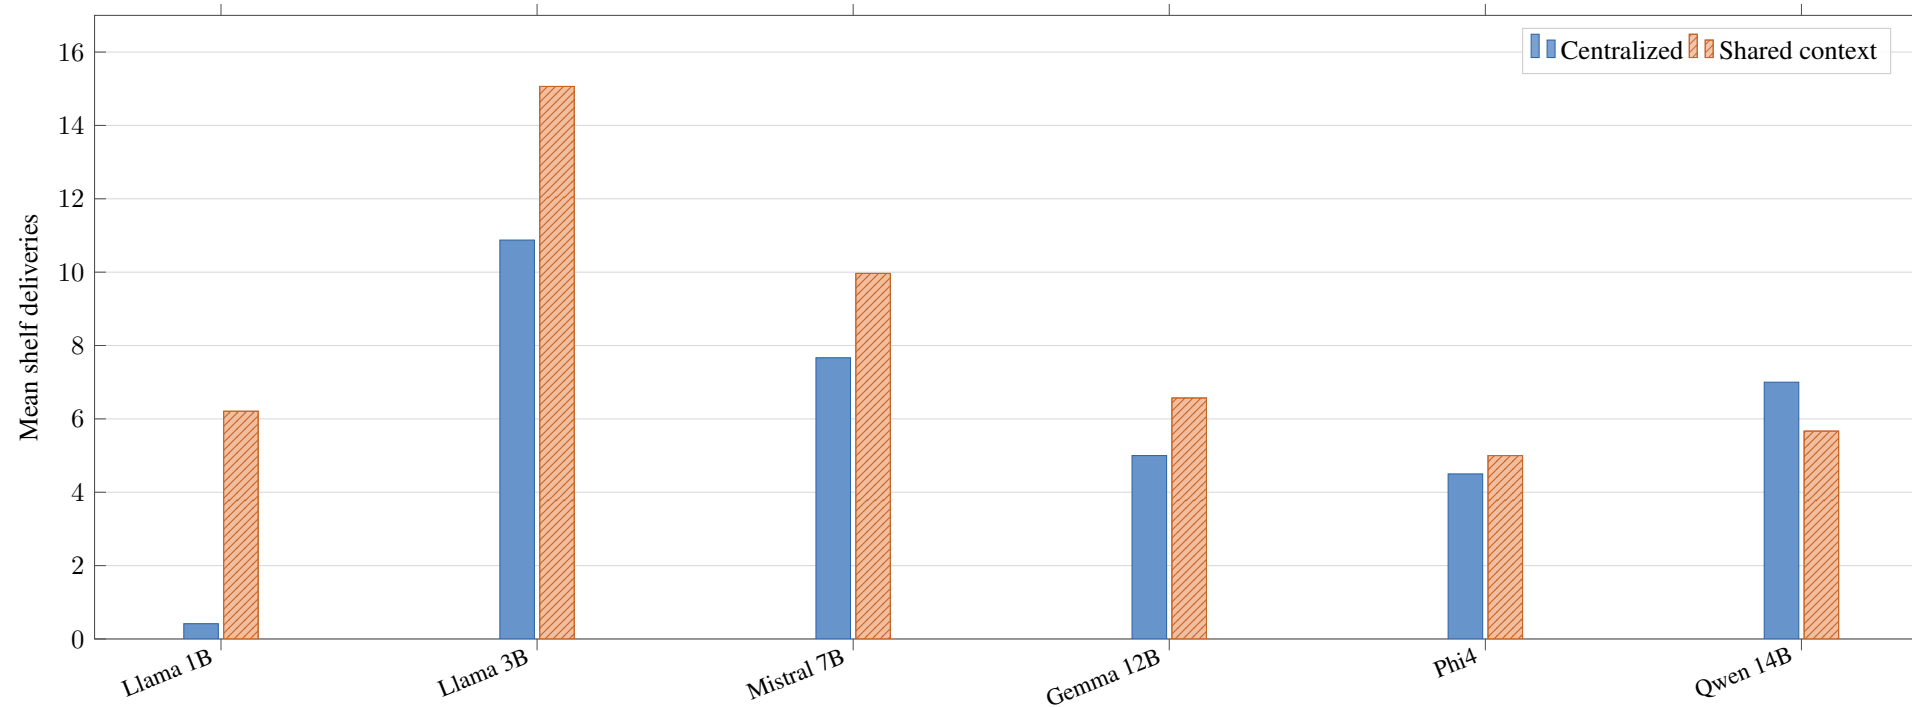
The centralized and shared-context groups each show the JSON and natural-language configuration side by side.

Mean delivery count by planning architecture and scenario



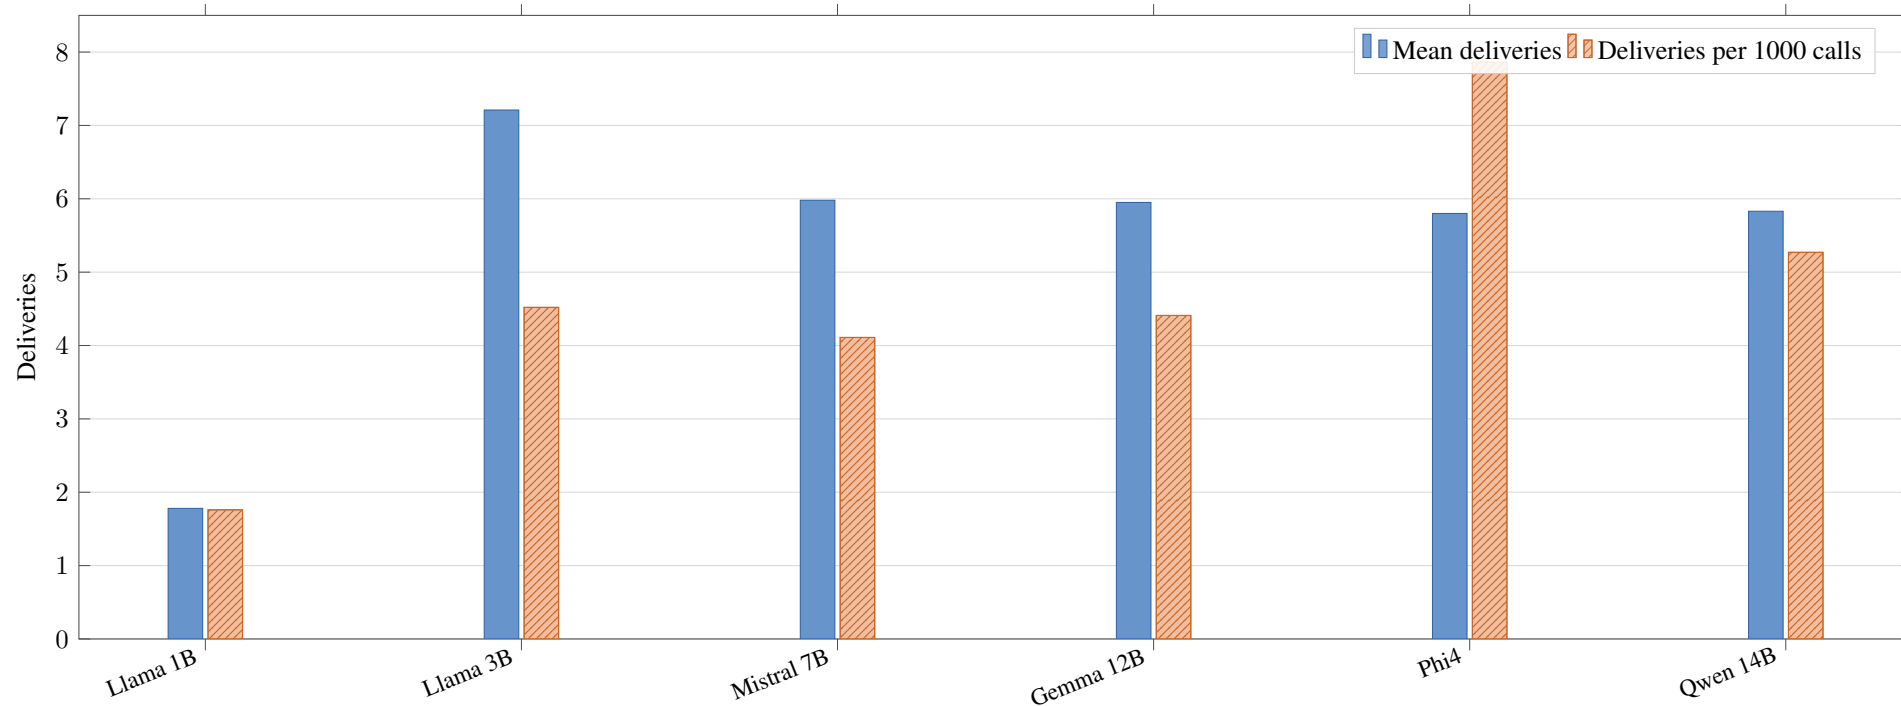
Grouped bars retain the ordered scenario presentation while treating every scenario as a discrete category. Each bar summarizes 16 matched configuration-level observations.

Mean delivery count by planning architecture and model



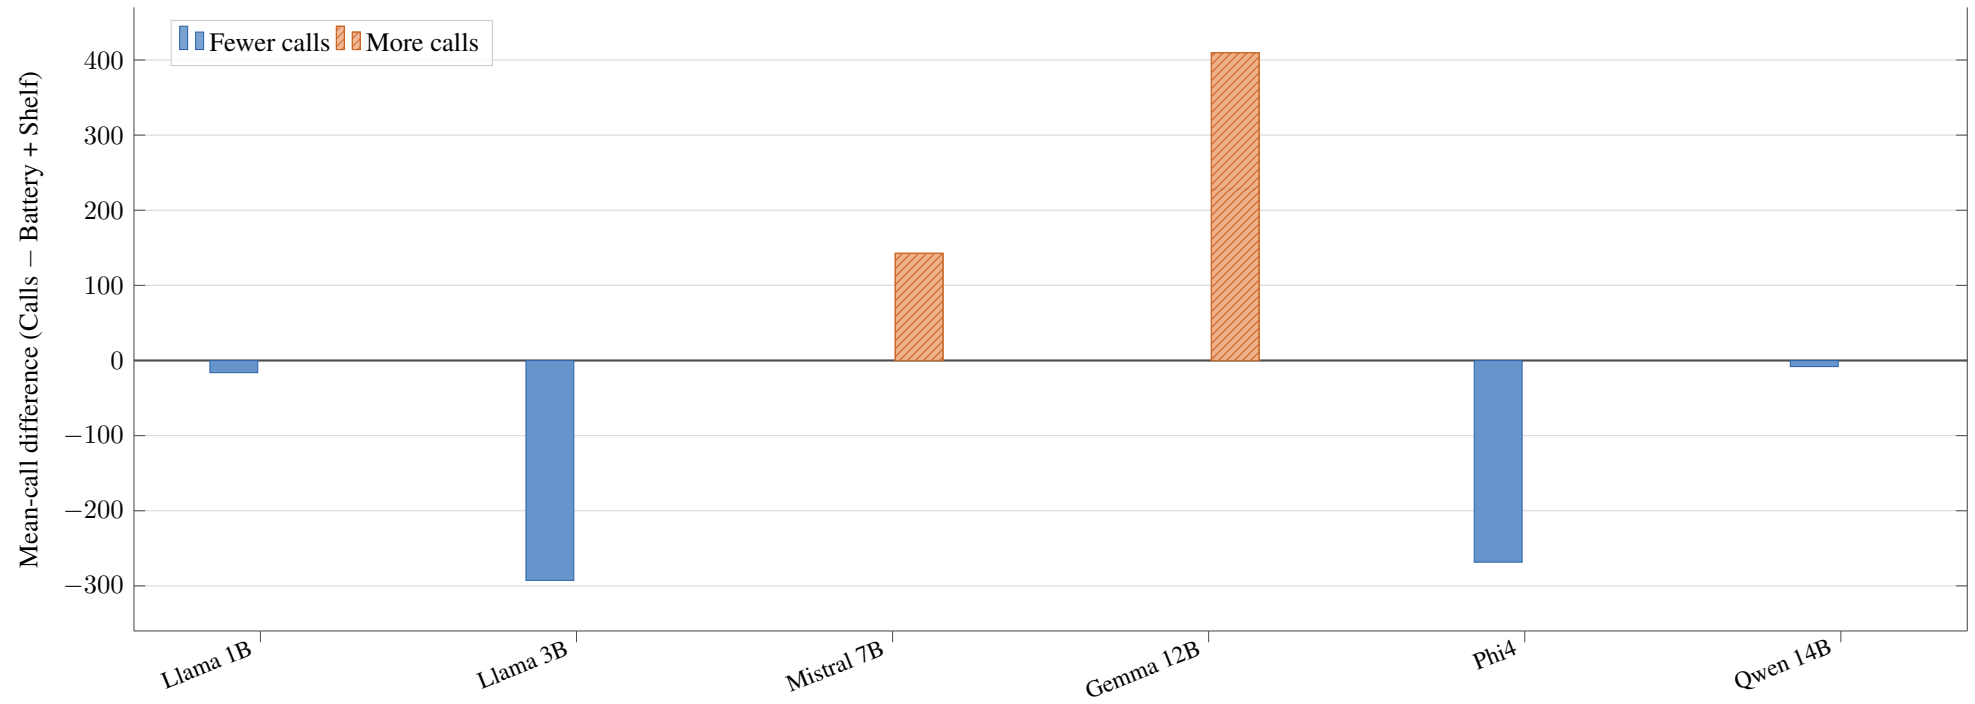
Models are ordered by nominal parameter size and remain discrete categories. Each bar summarizes 12 matched configuration-level observations.

Matched delivery performance by model



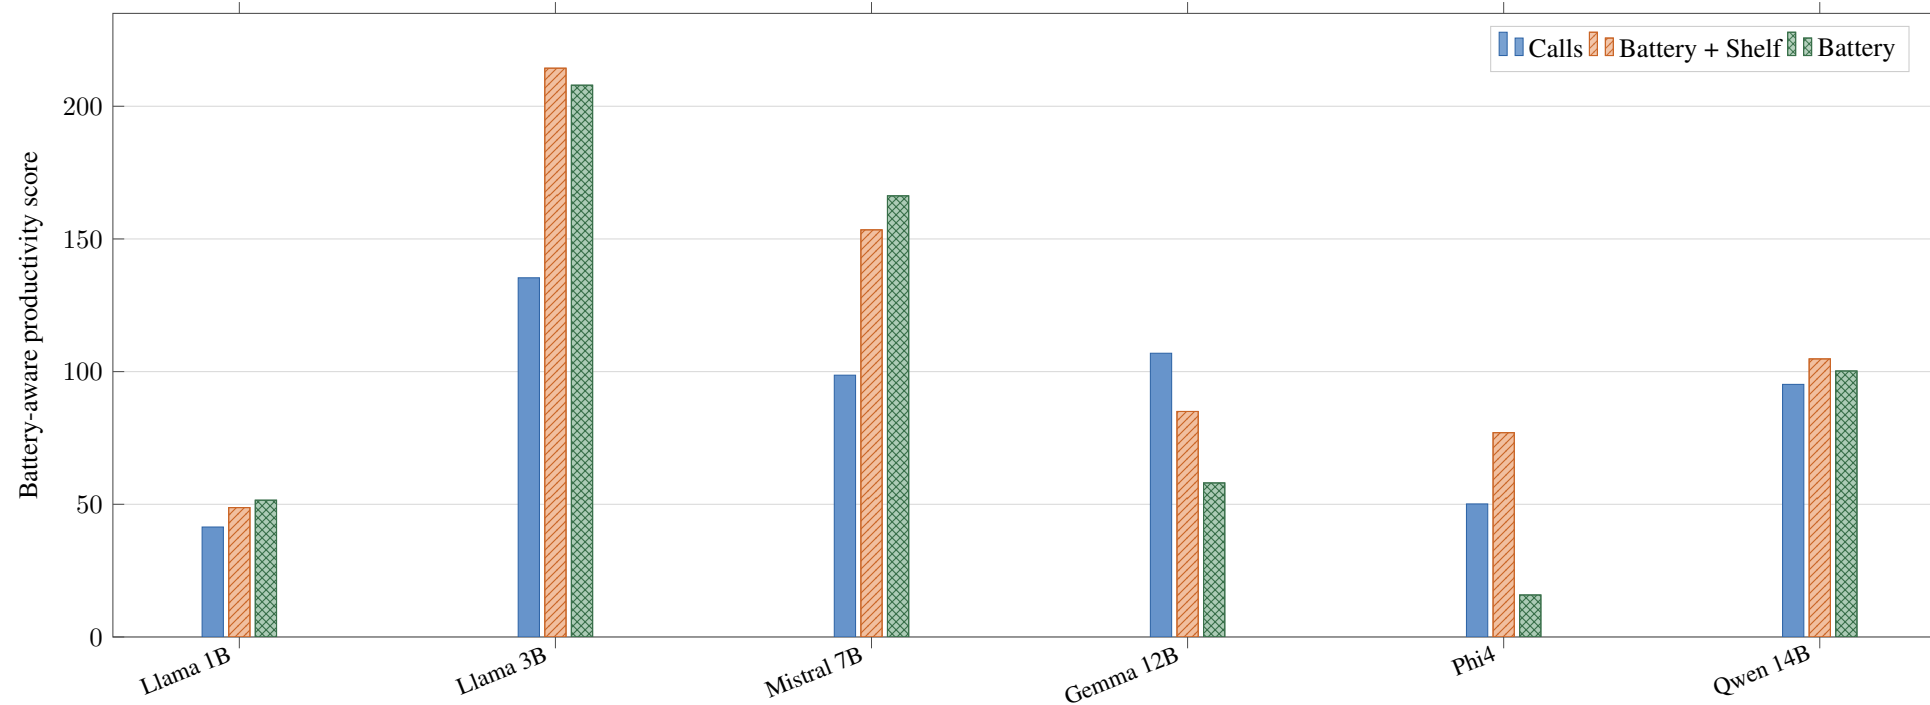
Mean delivery count and deliveries per 1000 LLM calls across the same 48 matched configurations for every model.

Change in LLM calls under the call-reduction objective



Difference is Calls minus Battery + Shelf. Negative bars indicate fewer calls under the call-reduction objective; positive bars indicate more calls.

Battery-aware productivity by objective and model



Descriptive battery-aware productivity score for the three matched objective configurations.